

# softmax

February 25, 2019

## 1 Softmax exercise

*Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission.*

In this exercise, you will:

- implement a fully-vectorized **loss function** for the Softmax classifier
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** with numerical gradient
- use a validation set to **tune the learning rate and regularization** strength
- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

```
In [1]: import random
import numpy as np
from data_utils import load_CIFAR10
import matplotlib.pyplot as plt

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

In [2]: def rel_error(out, correct_out):
    return np.sum(abs(out - correct_out) / (abs(out) + abs(correct_out)))

In [3]: def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000, num_dev=500):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for the linear classifier. These are the same steps as we used for the
    Softmax, but condensed to a single function.
    """
    # Load the raw CIFAR-10 data
```

```

cifar10_dir = 'datasets/cifar-10-batches-py'
X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

# subsample the data
mask = range(num_training, num_training + num_validation)
X_val = X_train[mask]
y_val = y_train[mask]
mask = range(num_training)
X_train = X_train[mask]
y_train = y_train[mask]
mask = range(num_test)
X_test = X_test[mask]
y_test = y_test[mask]

# # We will also make a development set, which is a small subset of
# the training set.
mask = np.random.choice(num_training, num_dev, replace=False)
X_dev = X_train[mask]
y_dev = y_train[mask]

# Preprocessing: reshape the image data into rows
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_val = np.reshape(X_val, (X_val.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))

# Normalize the data: subtract the mean image
mean_image = np.mean(X_train, axis = 0)
X_train -= mean_image
X_val -= mean_image
X_test -= mean_image
X_dev -= mean_image

# add bias dimension and transform into columns
X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])

return X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev = get_CIFAR10_data()
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)

```

```

print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
print('dev data shape: ', X_dev.shape)
print('dev labels shape: ', y_dev.shape)

```

```

Train data shape: (49000, 3073)
Train labels shape: (49000,)
Validation data shape: (1000, 3073)
Validation labels shape: (1000,)
Test data shape: (1000, 3073)
Test labels shape: (1000,)
dev data shape: (500, 3073)
dev labels shape: (500,)

```

```

In [4]: # Create one-hot vectors for label
num_class = 10
y_train_oh = np.zeros((y_train.shape[0], 10))
y_train_oh[np.arange(y_train.shape[0]), y_train] = 1
y_val_oh = np.zeros((y_val.shape[0], 10))
y_val_oh[np.arange(y_val.shape[0]), y_val] = 1
y_test_oh = np.zeros((y_test.shape[0], 10))
y_test_oh[np.arange(y_test.shape[0]), y_test] = 1

y_dev_oh = np.zeros((y_dev.shape[0], 10))
y_dev_oh[np.arange(y_dev.shape[0]), y_dev] = 1

```

## 2 Regression as classifier

The most simple and straightforward approach to learn a classifier is to map the input data (raw image values) to class label (one-hot vector). The loss function is defined as following:

$$\mathcal{L} = \frac{1}{n} \|\mathbf{XW} - \mathbf{y}\|_F^2 \quad (1)$$

Where: \*  $\mathbf{W} \in \mathbb{R}^{(d+1) \times C}$ : Classifier weight \*  $\mathbf{X} \in \mathbb{R}^{n \times (d+1)}$ : Dataset \*  $\mathbf{y} \in \mathbb{R}^{n \times C}$ : Class label (one-hot vector)

## 3 Optimization

Given the loss function (1), the next problem is how to solve the weight  $\mathbf{W}$ . We now discuss 2 approaches: \* Random search \* Closed-form solution

### 3.1 Random search

```

In [5]: bestloss = float('inf')
        for num in range(100):
            W = np.random.randn(3073, 10) * 0.0001

```

```

loss = np.linalg.norm(X_dev.dot(W) - y_dev_oh)
if (loss < bestloss):
    bestloss = loss
    bestW = W
print('in attempt %d the loss was %f, best %f' % (num, loss, bestloss))

```

```

in attempt 0 the loss was 31.097852, best 31.097852
in attempt 1 the loss was 35.986907, best 31.097852
in attempt 2 the loss was 32.936750, best 31.097852
in attempt 3 the loss was 32.263297, best 31.097852
in attempt 4 the loss was 33.316720, best 31.097852
in attempt 5 the loss was 33.530687, best 31.097852
in attempt 6 the loss was 33.089863, best 31.097852
in attempt 7 the loss was 33.797979, best 31.097852
in attempt 8 the loss was 31.849634, best 31.097852
in attempt 9 the loss was 31.078356, best 31.078356
in attempt 10 the loss was 32.901067, best 31.078356
in attempt 11 the loss was 32.724453, best 31.078356
in attempt 12 the loss was 34.693179, best 31.078356
in attempt 13 the loss was 33.269654, best 31.078356
in attempt 14 the loss was 33.896164, best 31.078356
in attempt 15 the loss was 31.727043, best 31.078356
in attempt 16 the loss was 31.291954, best 31.078356
in attempt 17 the loss was 34.024641, best 31.078356
in attempt 18 the loss was 31.856058, best 31.078356
in attempt 19 the loss was 34.706234, best 31.078356
in attempt 20 the loss was 32.552190, best 31.078356
in attempt 21 the loss was 30.909683, best 30.909683
in attempt 22 the loss was 33.826905, best 30.909683
in attempt 23 the loss was 31.748100, best 30.909683
in attempt 24 the loss was 33.683554, best 30.909683
in attempt 25 the loss was 33.235776, best 30.909683
in attempt 26 the loss was 33.766795, best 30.909683
in attempt 27 the loss was 32.458615, best 30.909683
in attempt 28 the loss was 32.793571, best 30.909683
in attempt 29 the loss was 32.980274, best 30.909683
in attempt 30 the loss was 33.646236, best 30.909683
in attempt 31 the loss was 32.000726, best 30.909683
in attempt 32 the loss was 34.343205, best 30.909683
in attempt 33 the loss was 32.809490, best 30.909683
in attempt 34 the loss was 30.753257, best 30.753257
in attempt 35 the loss was 33.363264, best 30.753257
in attempt 36 the loss was 32.521487, best 30.753257
in attempt 37 the loss was 32.286729, best 30.753257
in attempt 38 the loss was 33.466272, best 30.753257
in attempt 39 the loss was 33.767820, best 30.753257
in attempt 40 the loss was 34.014519, best 30.753257
in attempt 41 the loss was 33.604679, best 30.753257

```

in attempt 42 the loss was 33.021665, best 30.753257  
in attempt 43 the loss was 34.238782, best 30.753257  
in attempt 44 the loss was 31.868839, best 30.753257  
in attempt 45 the loss was 34.560101, best 30.753257  
in attempt 46 the loss was 33.859088, best 30.753257  
in attempt 47 the loss was 33.872607, best 30.753257  
in attempt 48 the loss was 35.401592, best 30.753257  
in attempt 49 the loss was 32.297212, best 30.753257  
in attempt 50 the loss was 33.731449, best 30.753257  
in attempt 51 the loss was 32.974246, best 30.753257  
in attempt 52 the loss was 33.360920, best 30.753257  
in attempt 53 the loss was 36.814551, best 30.753257  
in attempt 54 the loss was 32.383545, best 30.753257  
in attempt 55 the loss was 33.313233, best 30.753257  
in attempt 56 the loss was 32.943898, best 30.753257  
in attempt 57 the loss was 35.537028, best 30.753257  
in attempt 58 the loss was 31.539891, best 30.753257  
in attempt 59 the loss was 31.464733, best 30.753257  
in attempt 60 the loss was 32.306981, best 30.753257  
in attempt 61 the loss was 31.490039, best 30.753257  
in attempt 62 the loss was 32.139464, best 30.753257  
in attempt 63 the loss was 33.191718, best 30.753257  
in attempt 64 the loss was 32.380865, best 30.753257  
in attempt 65 the loss was 32.790204, best 30.753257  
in attempt 66 the loss was 30.227760, best 30.227760  
in attempt 67 the loss was 30.776805, best 30.227760  
in attempt 68 the loss was 32.365543, best 30.227760  
in attempt 69 the loss was 36.965447, best 30.227760  
in attempt 70 the loss was 32.575221, best 30.227760  
in attempt 71 the loss was 30.665061, best 30.227760  
in attempt 72 the loss was 33.697618, best 30.227760  
in attempt 73 the loss was 35.596273, best 30.227760  
in attempt 74 the loss was 32.185674, best 30.227760  
in attempt 75 the loss was 32.197692, best 30.227760  
in attempt 76 the loss was 32.185795, best 30.227760  
in attempt 77 the loss was 33.575504, best 30.227760  
in attempt 78 the loss was 32.962141, best 30.227760  
in attempt 79 the loss was 34.287886, best 30.227760  
in attempt 80 the loss was 32.794918, best 30.227760  
in attempt 81 the loss was 33.452308, best 30.227760  
in attempt 82 the loss was 33.735422, best 30.227760  
in attempt 83 the loss was 32.112291, best 30.227760  
in attempt 84 the loss was 34.023853, best 30.227760  
in attempt 85 the loss was 32.760048, best 30.227760  
in attempt 86 the loss was 32.747929, best 30.227760  
in attempt 87 the loss was 32.216412, best 30.227760  
in attempt 88 the loss was 33.294907, best 30.227760  
in attempt 89 the loss was 31.712067, best 30.227760

```

in attempt 90 the loss was 31.398513, best 30.227760
in attempt 91 the loss was 34.414248, best 30.227760
in attempt 92 the loss was 32.986523, best 30.227760
in attempt 93 the loss was 35.948089, best 30.227760
in attempt 94 the loss was 33.894141, best 30.227760
in attempt 95 the loss was 32.095476, best 30.227760
in attempt 96 the loss was 32.895736, best 30.227760
in attempt 97 the loss was 32.178781, best 30.227760
in attempt 98 the loss was 33.690184, best 30.227760
in attempt 99 the loss was 32.605644, best 30.227760

```

In [6]: # How bestW perform:

```

print('Accuracy on train set: ', np.sum(np.argmax(np.abs(1 - X_dev.dot(bestW)), axis=1) == y_train), axis=0)
print('Accuracy on test set: ', np.sum(np.argmax(np.abs(1 - X_test.dot(bestW)), axis=1) == y_test), axis=0)

```

Accuracy on train set: 15.8

Accuracy on test set: 16.900000000000002

You can clearly see that the performance is very low, almost at the random level.

## 3.2 Closed-form solution

The closed-form solution is achieved by:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}} = \frac{2}{n} \mathbf{X}^T (\mathbf{XW} - \mathbf{y}) = 0$$

$$\Leftrightarrow \mathbf{W}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

In [7]: #####  
# TODO: #

# Implement the closed-form solution of the weight W. #

#####

X\_train\_trans = np.transpose(X\_train)

pre = np.matmul(X\_train\_trans, X\_train)

pre\_inverse = np.linalg.inv(pre)

pre\_pre = np.matmul(pre\_inverse, X\_train\_trans)

W = np.matmul(pre\_pre, y\_train\_oh)

#####

# END OF YOUR CODE #

#####

In [8]: # Check accuracy:

```

print('Train set accuracy: ', np.sum(np.argmax(np.abs(1 - X_train.dot(W)), axis=1) == y_train), axis=0)
print('Test set accuracy: ', np.sum(np.argmax(np.abs(1 - X_test.dot(W)), axis=1) == y_test), axis=0)

```

Train set accuracy: 51.163265306122454  
 Test set accuracy: 36.199999999999996

Now, you can see that the performance is much better.

### 3.3 Regularization

A simple way to improve performance is to include the L2-regularization penalty.

$$\mathcal{L} = \frac{1}{n} \|\mathbf{X}\mathbf{W} - \mathbf{y}\|_F^2 + \lambda \|\mathbf{W}\|_F^2 \quad (2)$$

The closed-form solution now is:

$$\Leftrightarrow \mathbf{W}^* = (\mathbf{X}^T \mathbf{X} + \lambda n \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

```
In [9]: # try several values of lambda to see how it helps:
        lambdas = [0.01, 0.1, 1, 10, 100, 1000, 10000, 100000]
        train_acc = np.zeros((len(lambdas)))
        test_acc = np.zeros((len(lambdas)))
        for i in range(len(lambdas)):
            l = lambdas[i]
            n,d = X_train.shape[0], X_train.shape[1]
            #####
            # TODO:                                     #
            # Implement the closed-form solution of the weight W with regularization. #
            #####
            X_train_trans = np.transpose(X_train)
            pre_1 = np.matmul(X_train_trans, X_train)
            pre = pre_1 + l * n * np.identity(pre_1.shape[0])
            pre_inverse = np.linalg.inv(pre)
            pre_pre = np.matmul(pre_inverse, X_train_trans)
            W = np.matmul(pre_pre, y_train_oh)

            #####
            #                                     END OF YOUR CODE                                     #
            #####
            train_acc[i] = np.sum(np.argmax(np.abs(1 - X_train.dot(W)), axis=1) == y_train).astype(np.float)
            print('Train set accuracy of ', l, ' : ', train_acc[i])
            test_acc[i] = np.sum(np.argmax(np.abs(1 - X_test.dot(W)), axis=1) == y_test).astype(np.float)
            print('Test set accuracy of ', l, ' : ', test_acc[i])
```

```
Train set accuracy of 0.01 : 51.159183673469386
Test set accuracy of 0.01 : 36.3
Train set accuracy of 0.1 : 51.069387755102035
Test set accuracy of 0.1 : 36.3
Train set accuracy of 1 : 50.39591836734694
Test set accuracy of 1 : 37.1
```

```

Train set accuracy of 10 : 48.71632653061225
Test set accuracy of 10 : 37.8
Train set accuracy of 100 : 46.49183673469388
Test set accuracy of 100 : 39.2
Train set accuracy of 1000 : 44.304081632653066
Test set accuracy of 1000 : 40.5
Train set accuracy of 10000 : 41.55714285714286
Test set accuracy of 10000 : 40.2
Train set accuracy of 100000 : 37.30204081632653
Test set accuracy of 100000 : 38.1


```

```

In [10]: plt.semilogx(lambdas, train_acc, 'r', label="Training accuracy")
         plt.semilogx(lambdas, test_acc, 'g', label="Testing accuracy")

plt.legend()
plt.grid(True)
plt.show()

```




**Question:** Try to explain why the performances on the training and test set have such behaviors as we change the value of  $\lambda$ .

**Your answer:** When  $\lambda$  is small,  $W$  tends to be overfitting. When  $\lambda$  is 1000, it has the most optimized result for testing set.

### 3.4 Softmax Classifier

The predicted probability for the  $j$ -th class given a sample vector  $x$  and a weight  $W$  is:





$$P(y = j | x) = \frac{e^{-xw_j}}{\sum_{c=1}^C e^{-xw_c}}$$

Your code for this section will all be written inside **classifiers/softmax.py**.

```
In [14]: # First implement the naive softmax loss function with nested loops.
# Open the file classifiers/softmax.py and implement the
# softmax_loss_naive function.

from classifiers.softmax import softmax_loss_naive
import time

# Generate a random softmax weight matrix and use it to compute the loss.
W = np.random.randn(3073, 10) * 0.0001
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

# As a rough sanity check, our loss should be something close to -log(0.1).
print('loss: %f' % loss)
print('sanity check: %f' % (-np.log(0.1)))

loss: 2.332788
sanity check: 2.302585
```

**Question:** Why do we expect our loss to be close to  $-\log(0.1)$ ? Explain briefly.\*\*

**Your answer:** Because we have in total 10 classes. So we compare our result with random guessing.

## 4 Optimization

### 4.1 Random search

```
In [15]: bestloss = float('inf')
        for num in range(100):
            W = np.random.randn(3073, 10) * 0.0001
            loss, _ = softmax_loss_naive(W, X_dev, y_dev, 0.0)
            if (loss < bestloss):
                bestloss = loss
                bestW = W
            print('in attempt %d the loss was %f, best %f' % (num, loss, bestloss))
```

```
in attempt 0 the loss was 2.315287, best 2.315287
in attempt 1 the loss was 2.315728, best 2.315287
in attempt 2 the loss was 2.279245, best 2.279245
in attempt 3 the loss was 2.338965, best 2.279245
in attempt 4 the loss was 2.364009, best 2.279245
in attempt 5 the loss was 2.399972, best 2.279245
in attempt 6 the loss was 2.364766, best 2.279245
in attempt 7 the loss was 2.354514, best 2.279245
in attempt 8 the loss was 2.409205, best 2.279245
in attempt 9 the loss was 2.323656, best 2.279245
in attempt 10 the loss was 2.397302, best 2.279245
in attempt 11 the loss was 2.339812, best 2.279245
in attempt 12 the loss was 2.379946, best 2.279245
in attempt 13 the loss was 2.315949, best 2.279245
in attempt 14 the loss was 2.361623, best 2.279245
in attempt 15 the loss was 2.318217, best 2.279245
in attempt 16 the loss was 2.368780, best 2.279245
in attempt 17 the loss was 2.363134, best 2.279245
in attempt 18 the loss was 2.366291, best 2.279245
in attempt 19 the loss was 2.339898, best 2.279245
in attempt 20 the loss was 2.355486, best 2.279245
in attempt 21 the loss was 2.334065, best 2.279245
in attempt 22 the loss was 2.315165, best 2.279245
in attempt 23 the loss was 2.351390, best 2.279245
in attempt 24 the loss was 2.364131, best 2.279245
in attempt 25 the loss was 2.384957, best 2.279245
in attempt 26 the loss was 2.358619, best 2.279245
in attempt 27 the loss was 2.376561, best 2.279245
in attempt 28 the loss was 2.329522, best 2.279245
in attempt 29 the loss was 2.337146, best 2.279245
in attempt 30 the loss was 2.353836, best 2.279245
```

in attempt 31 the loss was 2.321441, best 2.279245  
in attempt 32 the loss was 2.333559, best 2.279245  
in attempt 33 the loss was 2.330200, best 2.279245  
in attempt 34 the loss was 2.307366, best 2.279245  
in attempt 35 the loss was 2.361179, best 2.279245  
in attempt 36 the loss was 2.323411, best 2.279245  
in attempt 37 the loss was 2.337375, best 2.279245  
in attempt 38 the loss was 2.381666, best 2.279245  
in attempt 39 the loss was 2.351327, best 2.279245  
in attempt 40 the loss was 2.355621, best 2.279245  
in attempt 41 the loss was 2.401918, best 2.279245  
in attempt 42 the loss was 2.371483, best 2.279245  
in attempt 43 the loss was 2.397893, best 2.279245  
in attempt 44 the loss was 2.336500, best 2.279245  
in attempt 45 the loss was 2.306861, best 2.279245  
in attempt 46 the loss was 2.343061, best 2.279245  
in attempt 47 the loss was 2.362867, best 2.279245  
in attempt 48 the loss was 2.391230, best 2.279245  
in attempt 49 the loss was 2.374461, best 2.279245  
in attempt 50 the loss was 2.383866, best 2.279245  
in attempt 51 the loss was 2.369601, best 2.279245  
in attempt 52 the loss was 2.399752, best 2.279245  
in attempt 53 the loss was 2.346172, best 2.279245  
in attempt 54 the loss was 2.309303, best 2.279245  
in attempt 55 the loss was 2.366009, best 2.279245  
in attempt 56 the loss was 2.312034, best 2.279245  
in attempt 57 the loss was 2.399871, best 2.279245  
in attempt 58 the loss was 2.369382, best 2.279245  
in attempt 59 the loss was 2.353133, best 2.279245  
in attempt 60 the loss was 2.318029, best 2.279245  
in attempt 61 the loss was 2.375471, best 2.279245  
in attempt 62 the loss was 2.335294, best 2.279245  
in attempt 63 the loss was 2.388040, best 2.279245  
in attempt 64 the loss was 2.296572, best 2.279245  
in attempt 65 the loss was 2.327079, best 2.279245  
in attempt 66 the loss was 2.376346, best 2.279245  
in attempt 67 the loss was 2.371548, best 2.279245  
in attempt 68 the loss was 2.328372, best 2.279245  
in attempt 69 the loss was 2.379915, best 2.279245  
in attempt 70 the loss was 2.377497, best 2.279245  
in attempt 71 the loss was 2.297419, best 2.279245  
in attempt 72 the loss was 2.311660, best 2.279245  
in attempt 73 the loss was 2.366952, best 2.279245  
in attempt 74 the loss was 2.342677, best 2.279245  
in attempt 75 the loss was 2.351796, best 2.279245  
in attempt 76 the loss was 2.318753, best 2.279245  
in attempt 77 the loss was 2.342505, best 2.279245  
in attempt 78 the loss was 2.333173, best 2.279245

```

in attempt 79 the loss was 2.351038, best 2.279245
in attempt 80 the loss was 2.388554, best 2.279245
in attempt 81 the loss was 2.326636, best 2.279245
in attempt 82 the loss was 2.343161, best 2.279245
in attempt 83 the loss was 2.345845, best 2.279245
in attempt 84 the loss was 2.378235, best 2.279245
in attempt 85 the loss was 2.323155, best 2.279245
in attempt 86 the loss was 2.418805, best 2.279245
in attempt 87 the loss was 2.382597, best 2.279245
in attempt 88 the loss was 2.369079, best 2.279245
in attempt 89 the loss was 2.340579, best 2.279245
in attempt 90 the loss was 2.364311, best 2.279245
in attempt 91 the loss was 2.331434, best 2.279245
in attempt 92 the loss was 2.293285, best 2.279245
in attempt 93 the loss was 2.394173, best 2.279245
in attempt 94 the loss was 2.362335, best 2.279245
in attempt 95 the loss was 2.370387, best 2.279245
in attempt 96 the loss was 2.323099, best 2.279245
in attempt 97 the loss was 2.334008, best 2.279245
in attempt 98 the loss was 2.317522, best 2.279245
in attempt 99 the loss was 2.355111, best 2.279245

```

```

In [16]: # How bestW perform on trainset
         scores = X_train.dot(bestW)
         y_pred = np.argmax(scores, axis=1)
         print('Accuracy on train set %f' % np.mean(y_pred == y_train))

         # evaluate performance of test set
         scores = X_test.dot(bestW)
         y_pred = np.argmax(scores, axis=1)
         print('Accuracy on test set %f' % np.mean(y_pred == y_test))

```

Accuracy on train set 0.131245

Accuracy on test set 0.124000

Compare the performance when using random search with *regression classifier* and *softmax classifier*. You can see how much useful the softmax classifier is.

## 4.2 Stochastic Gradient descent

Even though it is possible to achieve closed-form solution with softmax classifier, it would be more complicated. In fact, we could achieve very good results with gradient descent approach. Additionally, in case of very large dataset, it is impossible to load the whole dataset into the memory. Gradient descent can help to optimize the loss function in batch.

$$\mathbf{W}^{t+1} = \mathbf{W}^t - \alpha \frac{\partial \mathcal{L}(\mathbf{x}; \mathbf{W}^t)}{\partial \mathbf{W}^t}$$

Where  $\alpha$  is the learning rate,  $\mathcal{L}$  is a loss function, and  $\mathbf{x}$  is a batch of training dataset.

```
In [17]: # Complete the implementation of softmax_loss_naive and implement a (naive)
# version of the gradient that uses nested loops.
```

```
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)
```

```
# Use numeric gradient checking as a debugging tool.
```

```
# The numeric gradient should be close to the analytic gradient.
```

```
from gradient_check import grad_check_sparse
```

```
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 0.0)[0]
```

```
grad_numerical = grad_check_sparse(f, W, grad, 10)
```

```
# gradient check with regularization
```

```
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 1e2)
```

```
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 1e2)[0]
```

```
grad_numerical = grad_check_sparse(f, W, grad, 10)
```

```
numerical: -0.863772 analytic: -0.863772, relative error: 3.286460e-08
```

```
numerical: -0.830557 analytic: -0.830557, relative error: 9.243450e-08
```

```
numerical: 2.159302 analytic: 2.159302, relative error: 3.404888e-08
```

```
numerical: -0.075499 analytic: -0.075499, relative error: 2.376996e-07
```

```
numerical: -5.666792 analytic: -5.666792, relative error: 2.276835e-09
```

```
numerical: -2.117591 analytic: -2.117591, relative error: 5.739073e-09
```

```
numerical: -0.995018 analytic: -0.995018, relative error: 1.605112e-08
```

```
numerical: 0.199430 analytic: 0.199430, relative error: 1.770890e-07
```

```
numerical: -0.060309 analytic: -0.060309, relative error: 9.127597e-08
```

```
numerical: -2.069928 analytic: -2.069928, relative error: 1.044419e-08
```

```
numerical: 0.176114 analytic: 0.176114, relative error: 2.150958e-07
```

```
numerical: -4.152779 analytic: -4.152779, relative error: 4.241791e-09
```

```
numerical: 0.782594 analytic: 0.782594, relative error: 1.654405e-08
```

```
numerical: 0.193260 analytic: 0.193260, relative error: 1.609427e-07
```

```
numerical: -2.683427 analytic: -2.683427, relative error: 1.281667e-09
```

```
numerical: -0.846214 analytic: -0.846214, relative error: 1.180543e-08
```

```
numerical: -0.289115 analytic: -0.289115, relative error: 2.914200e-08
```

```
numerical: 0.035788 analytic: 0.035788, relative error: 1.307111e-06
```

```
numerical: -3.161925 analytic: -3.161925, relative error: 5.695438e-09
```

```
numerical: 0.322899 analytic: 0.322899, relative error: 1.329794e-07
```

```
In [24]: # Now that we have a naive implementation of the softmax loss function and its gradient,
# implement a vectorized version in softmax_loss_vectorized.
```

```
# The two versions should compute the same results, but the vectorized version should
# much faster.
```

```
tic = time.time()
```

```
loss_naive, grad_naive = softmax_loss_naive(W, X_dev, y_dev, 0.00001)
```

```
toc = time.time()
```

```
print('naive loss: %e computed in %fs' % (loss_naive, toc - tic))
```

```

from classifiers.softmax import softmax_loss_vectorized
tic = time.time()
loss_vectorized, grad_vectorized = softmax_loss_vectorized(W, X_dev, y_dev, 0.00001)
toc = time.time()
print('vectorized loss: %e computed in %fs' % (loss_vectorized, toc - tic))

# We use the Frobenius norm to compare the two versions
# of the gradient.
grad_difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')
print('Loss difference: %f' % np.abs(loss_naive - loss_vectorized))
print('Gradient difference: %f' % grad_difference)

```

naive loss: 2.355111e+00 computed in 11.699731s  
vectorized loss: 8.587773e+00 computed in 0.009061s  
Loss difference: 6.232662  
Gradient difference: 141.386003

```

In [27]: from classifiers.linear_classifier import *
         # from classifiers import linear_classifier

         classifier = Softmax()
         tic = time.time()
         loss_hist = classifier.train(X_train, y_train, learning_rate=1e-7, reg=5e4,
                                     num_iters=1500, verbose=True)

         toc = time.time()
         print('That took %fs' % (toc - tic))

```

```

iteration 0 / 1500: loss 780.933791
iteration 100 / 1500: loss 781.235478
iteration 200 / 1500: loss 781.834763
iteration 300 / 1500: loss 784.080822
iteration 400 / 1500: loss 785.783550
iteration 500 / 1500: loss 788.199196
iteration 600 / 1500: loss 792.060046
iteration 700 / 1500: loss 795.235677
iteration 800 / 1500: loss 799.411085
iteration 900 / 1500: loss 804.025054
iteration 1000 / 1500: loss 808.421515
iteration 1100 / 1500: loss 815.391185
iteration 1200 / 1500: loss 820.561515
iteration 1300 / 1500: loss 828.572252
iteration 1400 / 1500: loss 834.540254
That took 4.650405s

```

```

In [28]: # Write the Softmax.predict function and evaluate the performance on both the
         # training and validation set
         y_train_pred = classifier.predict(X_train)

```

```

print('training accuracy: %f' % (np.mean(y_train == y_train_pred), ))
y_val_pred = classifier.predict(X_val)
print('validation accuracy: %f' % (np.mean(y_val == y_val_pred), ))


```

training accuracy: 0.248061  
validation accuracy: 0.263000

```

In [29]: # A useful debugging strategy is to plot the loss as a function of
# iteration number:
plt.plot(loss_hist)
plt.xlabel('Iteration number')
plt.ylabel('Loss value')
plt.show()

```



```

In [30]: # evaluate on test set
# Evaluate the best softmax on test set
y_test_pred = classifier.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print('softmax on raw pixels final test set accuracy: %.2f' % (100*test_accuracy, ))

```

softmax on raw pixels final test set accuracy: 26.70

```

In [33]: # Use the validation set to tune hyperparameters (regularization strength and
# learning rate). You should experiment with different ranges for the learning

```

```

# rates and regularization strengths; if you are careful you should be able to
# get a classification accuracy of over 35% on the validation set.

import copy
best_val_acc = -1
best_softmax = None      # You may need to use copy.deepcopy(object)
learning_rates = [1e-4, 1e-5, 1e-6, 5e-7]
regularization_strengths = [1e-4, 1e-3, 1e-2, 1e-1, 1e0]
results = {}
#####
# TODO:                                     #
# Use the validation set to set the learning rate and regularization strength. #
# Save the best trained softmax classifier in best_softmax.                   #
#####
for lr in learning_rates:
    for rs in regularization_strengths:
        classifier = Softmax()
        classifier.train(X_train, y_train, learning_rate=lr, reg=rs,
                        num_iters=1500, verbose=True)
        y_val_pred = classifier.predict(X_val)
        val_acc = np.mean(y_val == y_val_pred)
        if val_acc > best_val_acc:
            best_softmax = classifier
            best_val_acc = val_acc

#####
#                                     END OF YOUR CODE                                     #
#####
print('best cross-validation accuracy: %.2f' % best_val_acc)

```

```
iteration 0 / 1500: loss 10.833682
```

```
/Users/lixingxuan/Desktop/Computer Vision/Week-01/forStudents/classifiers/softmax.py:90: RuntimeWarning: 
```

```
loss = np.sum(-np.log(correct))
```

```

iteration 100 / 1500: loss inf
iteration 200 / 1500: loss inf
iteration 300 / 1500: loss inf
iteration 400 / 1500: loss inf
iteration 500 / 1500: loss inf
iteration 600 / 1500: loss inf
iteration 700 / 1500: loss inf
iteration 800 / 1500: loss inf
iteration 900 / 1500: loss inf
iteration 1000 / 1500: loss inf

```



iteration 1100 / 1500: loss inf  
iteration 1200 / 1500: loss inf  
iteration 1300 / 1500: loss inf  
iteration 1400 / 1500: loss inf  
iteration 0 / 1500: loss 10.256135  
iteration 100 / 1500: loss inf  
iteration 200 / 1500: loss inf  
iteration 300 / 1500: loss inf  
iteration 400 / 1500: loss inf  
iteration 500 / 1500: loss inf  
iteration 600 / 1500: loss inf  
iteration 700 / 1500: loss inf  
iteration 800 / 1500: loss inf  
iteration 900 / 1500: loss inf  
iteration 1000 / 1500: loss inf  
iteration 1100 / 1500: loss inf  
iteration 1200 / 1500: loss inf  
iteration 1300 / 1500: loss inf  
iteration 1400 / 1500: loss inf  
iteration 0 / 1500: loss 11.253733  
iteration 100 / 1500: loss inf  
iteration 200 / 1500: loss inf  
iteration 300 / 1500: loss inf  
iteration 400 / 1500: loss inf  
iteration 500 / 1500: loss inf  
iteration 600 / 1500: loss inf  
iteration 700 / 1500: loss inf  
iteration 800 / 1500: loss inf  
iteration 900 / 1500: loss inf  
iteration 1000 / 1500: loss inf  
iteration 1100 / 1500: loss inf  
iteration 1200 / 1500: loss inf  
iteration 1300 / 1500: loss inf  
iteration 1400 / 1500: loss inf  
iteration 0 / 1500: loss 11.130422  
iteration 100 / 1500: loss inf  
iteration 200 / 1500: loss inf  
iteration 300 / 1500: loss inf  
iteration 400 / 1500: loss inf  
iteration 500 / 1500: loss inf  
iteration 600 / 1500: loss inf  
iteration 700 / 1500: loss inf  
iteration 800 / 1500: loss inf  
iteration 900 / 1500: loss inf  
iteration 1000 / 1500: loss inf  
iteration 1100 / 1500: loss inf  
iteration 1200 / 1500: loss inf  
iteration 1300 / 1500: loss inf

iteration 1400 / 1500: loss inf  
iteration 0 / 1500: loss 11.434804  
iteration 100 / 1500: loss inf  
iteration 200 / 1500: loss inf  
iteration 300 / 1500: loss inf  
iteration 400 / 1500: loss inf  
iteration 500 / 1500: loss inf  
iteration 600 / 1500: loss inf  
iteration 700 / 1500: loss inf  
iteration 800 / 1500: loss inf  
iteration 900 / 1500: loss inf  
iteration 1000 / 1500: loss inf  
iteration 1100 / 1500: loss inf  
iteration 1200 / 1500: loss inf  
iteration 1300 / 1500: loss inf  
iteration 1400 / 1500: loss inf  
iteration 0 / 1500: loss 10.616550  
iteration 100 / 1500: loss inf  
iteration 200 / 1500: loss inf  
iteration 300 / 1500: loss inf  
iteration 400 / 1500: loss inf  
iteration 500 / 1500: loss inf  
iteration 600 / 1500: loss inf  
iteration 700 / 1500: loss inf  
iteration 800 / 1500: loss inf  
iteration 900 / 1500: loss inf  
iteration 1000 / 1500: loss inf  
iteration 1100 / 1500: loss inf  
iteration 1200 / 1500: loss inf  
iteration 1300 / 1500: loss inf  
iteration 1400 / 1500: loss inf  
iteration 0 / 1500: loss 11.057112  
iteration 100 / 1500: loss inf  
iteration 200 / 1500: loss inf  
iteration 300 / 1500: loss inf  
iteration 400 / 1500: loss inf  
iteration 500 / 1500: loss inf  
iteration 600 / 1500: loss inf  
iteration 700 / 1500: loss inf  
iteration 800 / 1500: loss inf  
iteration 900 / 1500: loss inf  
iteration 1000 / 1500: loss inf  
iteration 1100 / 1500: loss inf  
iteration 1200 / 1500: loss inf  
iteration 1300 / 1500: loss inf  
iteration 1400 / 1500: loss inf  
iteration 0 / 1500: loss 10.969724  
iteration 100 / 1500: loss 100.659225

iteration 200 / 1500: loss inf  
iteration 300 / 1500: loss inf  
iteration 400 / 1500: loss inf  
iteration 500 / 1500: loss inf  
iteration 600 / 1500: loss inf  
iteration 700 / 1500: loss inf  
iteration 800 / 1500: loss inf  
iteration 900 / 1500: loss inf  
iteration 1000 / 1500: loss inf  
iteration 1100 / 1500: loss inf  
iteration 1200 / 1500: loss inf  
iteration 1300 / 1500: loss inf  
iteration 1400 / 1500: loss inf  
iteration 0 / 1500: loss 10.538474  
iteration 100 / 1500: loss inf  
iteration 200 / 1500: loss inf  
iteration 300 / 1500: loss inf  
iteration 400 / 1500: loss inf  
iteration 500 / 1500: loss inf  
iteration 600 / 1500: loss inf  
iteration 700 / 1500: loss inf  
iteration 800 / 1500: loss inf  
iteration 900 / 1500: loss inf  
iteration 1000 / 1500: loss inf  
iteration 1100 / 1500: loss inf  
iteration 1200 / 1500: loss inf  
iteration 1300 / 1500: loss inf  
iteration 1400 / 1500: loss inf  
iteration 0 / 1500: loss 11.880431  
iteration 100 / 1500: loss inf  
iteration 200 / 1500: loss inf  
iteration 300 / 1500: loss inf  
iteration 400 / 1500: loss inf  
iteration 500 / 1500: loss inf  
iteration 600 / 1500: loss inf  
iteration 700 / 1500: loss inf  
iteration 800 / 1500: loss inf  
iteration 900 / 1500: loss inf  
iteration 1000 / 1500: loss inf  
iteration 1100 / 1500: loss inf  
iteration 1200 / 1500: loss inf  
iteration 1300 / 1500: loss inf  
iteration 1400 / 1500: loss inf  
iteration 0 / 1500: loss 11.461731  
iteration 100 / 1500: loss 17.508315  
iteration 200 / 1500: loss 31.471587  
iteration 300 / 1500: loss 40.589631  
iteration 400 / 1500: loss 49.922190

iteration 500 / 1500: loss 73.070554  
iteration 600 / 1500: loss 74.045998  
iteration 700 / 1500: loss inf  
iteration 800 / 1500: loss 81.753876  
iteration 900 / 1500: loss inf  
iteration 1000 / 1500: loss 119.025187  
iteration 1100 / 1500: loss inf  
iteration 1200 / 1500: loss inf  
iteration 1300 / 1500: loss inf  
iteration 1400 / 1500: loss inf  
iteration 0 / 1500: loss 10.812399  
iteration 100 / 1500: loss 18.555991  
iteration 200 / 1500: loss 26.984466  
iteration 300 / 1500: loss 34.214327  
iteration 400 / 1500: loss 42.535130  
iteration 500 / 1500: loss 65.451920  
iteration 600 / 1500: loss 63.369286  
iteration 700 / 1500: loss inf  
iteration 800 / 1500: loss 82.947996  
iteration 900 / 1500: loss 102.974315  
iteration 1000 / 1500: loss inf  
iteration 1100 / 1500: loss inf  
iteration 1200 / 1500: loss inf  
iteration 1300 / 1500: loss inf  
iteration 1400 / 1500: loss inf  
iteration 0 / 1500: loss 10.637520  
iteration 100 / 1500: loss 17.337011  
iteration 200 / 1500: loss 25.298642  
iteration 300 / 1500: loss 38.467463  
iteration 400 / 1500: loss 54.939574  
iteration 500 / 1500: loss 60.414140  
iteration 600 / 1500: loss 72.455802  
iteration 700 / 1500: loss inf  
iteration 800 / 1500: loss 90.164946  
iteration 900 / 1500: loss 100.187124  
iteration 1000 / 1500: loss inf  
iteration 1100 / 1500: loss inf  
iteration 1200 / 1500: loss inf  
iteration 1300 / 1500: loss inf  
iteration 1400 / 1500: loss inf  
iteration 0 / 1500: loss 10.639573  
iteration 100 / 1500: loss 17.354602  
iteration 200 / 1500: loss 27.982728  
iteration 300 / 1500: loss 34.465859  
iteration 400 / 1500: loss 53.257845  
iteration 500 / 1500: loss 57.105250  
iteration 600 / 1500: loss 70.661566  
iteration 700 / 1500: loss 85.868442

iteration 800 / 1500: loss inf  
iteration 900 / 1500: loss 91.792896  
iteration 1000 / 1500: loss inf  
iteration 1100 / 1500: loss inf  
iteration 1200 / 1500: loss inf  
iteration 1300 / 1500: loss inf  
iteration 1400 / 1500: loss inf  
iteration 0 / 1500: loss 10.599205  
iteration 100 / 1500: loss 16.179100  
iteration 200 / 1500: loss 25.230719  
iteration 300 / 1500: loss 37.032863  
iteration 400 / 1500: loss 51.087448  
iteration 500 / 1500: loss 55.806179  
iteration 600 / 1500: loss 84.189182  
iteration 700 / 1500: loss inf  
iteration 800 / 1500: loss 98.865651  
iteration 900 / 1500: loss inf  
iteration 1000 / 1500: loss inf  
iteration 1100 / 1500: loss inf  
iteration 1200 / 1500: loss inf  
iteration 1300 / 1500: loss inf  
iteration 1400 / 1500: loss inf  
iteration 0 / 1500: loss 11.035841  
iteration 100 / 1500: loss 14.950523  
iteration 200 / 1500: loss 21.025990  
iteration 300 / 1500: loss 26.354624  
iteration 400 / 1500: loss 34.132120  
iteration 500 / 1500: loss 38.264403  
iteration 600 / 1500: loss 43.352790  
iteration 700 / 1500: loss 46.031743  
iteration 800 / 1500: loss 51.374813  
iteration 900 / 1500: loss 50.505969  
iteration 1000 / 1500: loss 61.946527  
iteration 1100 / 1500: loss 72.432727  
iteration 1200 / 1500: loss 81.072738  
iteration 1300 / 1500: loss 90.872571  
iteration 1400 / 1500: loss 71.133581  
iteration 0 / 1500: loss 10.704950  
iteration 100 / 1500: loss 13.705619  
iteration 200 / 1500: loss 17.180218  
iteration 300 / 1500: loss 22.420049  
iteration 400 / 1500: loss 28.359984  
iteration 500 / 1500: loss 34.363021  
iteration 600 / 1500: loss 40.114517  
iteration 700 / 1500: loss 44.531037  
iteration 800 / 1500: loss 52.417248  
iteration 900 / 1500: loss 57.963793  
iteration 1000 / 1500: loss 62.803546

iteration 1100 / 1500: loss 72.991386  
iteration 1200 / 1500: loss 79.633485  
iteration 1300 / 1500: loss 80.579326  
iteration 1400 / 1500: loss 93.773303  
iteration 0 / 1500: loss 10.588889  
iteration 100 / 1500: loss 14.113477  
iteration 200 / 1500: loss 16.690732  
iteration 300 / 1500: loss 20.885429  
iteration 400 / 1500: loss 27.943624  
iteration 500 / 1500: loss 35.677015  
iteration 600 / 1500: loss 41.977149  
iteration 700 / 1500: loss 43.387496  
iteration 800 / 1500: loss 56.239834  
iteration 900 / 1500: loss 53.695806  
iteration 1000 / 1500: loss 65.048353  
iteration 1100 / 1500: loss 56.544349  
iteration 1200 / 1500: loss 74.650788  
iteration 1300 / 1500: loss 67.251700  
iteration 1400 / 1500: loss 88.379485  
iteration 0 / 1500: loss 9.969727  
iteration 100 / 1500: loss 12.915702  
iteration 200 / 1500: loss 16.621113  
iteration 300 / 1500: loss 20.708916  
iteration 400 / 1500: loss 25.566202  
iteration 500 / 1500: loss 33.287885  
iteration 600 / 1500: loss 36.368469  
iteration 700 / 1500: loss 40.637098  
iteration 800 / 1500: loss 49.242602  
iteration 900 / 1500: loss 50.145033  
iteration 1000 / 1500: loss 51.928346  
iteration 1100 / 1500: loss 80.383780  
iteration 1200 / 1500: loss 72.144918  
iteration 1300 / 1500: loss 76.597436  
iteration 1400 / 1500: loss 69.619972  
iteration 0 / 1500: loss 10.484946  
iteration 100 / 1500: loss 12.304672  
iteration 200 / 1500: loss 19.542368  
iteration 300 / 1500: loss 20.984010  
iteration 400 / 1500: loss 29.869059  
iteration 500 / 1500: loss 34.094379  
iteration 600 / 1500: loss 36.014864  
iteration 700 / 1500: loss 46.991440  
iteration 800 / 1500: loss 49.816525  
iteration 900 / 1500: loss 51.478987  
iteration 1000 / 1500: loss 62.890280  
iteration 1100 / 1500: loss 66.512474  
iteration 1200 / 1500: loss 71.642227  
iteration 1300 / 1500: loss 69.214154

```
iteration 1400 / 1500: loss inf
best cross-validation accuracy: 0.27
```

```
In [34]: # evaluate on test set
         # Evaluate the best softmax on test set
         y_test_pred = best_softmax.predict(X_test)
         test_accuracy = np.mean(y_test == y_test_pred)
         print('softmax on raw pixels final test set accuracy: %.2f' % (100*test_accuracy, ))

softmax on raw pixels final test set accuracy: 26.40
```